





# Introduction to Operating Systems

## 1. Operating systems overview

Manuel – Fall 2019



# Hardware and software



A computer consists of:



- Hardware
- Software
  - Kernel mode
  - User mode

## Operating System (OS)

- Hardware is complicated to handle
- OS hides all the messy details
- OS manages resources for each program (time and space)
- Renders computer much easier to use

## The first days:

- Birth of modern computing: 19th century (Babbage)
- Vacuum tube: 1945–1955 (1st generation)
- Transistor: 1955–1965 (2nd generation)





Using the device:

- Program at most 40 steps
- Read input from cardboards
- Wire them on a plugboard
- Punch output on cardboards

Multiprogramming: 1965–1980 (3rd generation)

- Multiple jobs kept in memory at the same time
- CPU multiplexed among them

Multiprogramming: 1965–1980 (3rd generation)

- Multiple jobs kept in memory at the same time
- CPU multiplexed among them

Multiprogramming requires:

- Memory management: allocate memory to several jobs
- CPU scheduling: choose a job to be run
- Simultaneous Peripheral Operation On Line (SPOOL): load a new job from disk, run it, output it on disk



Most famous OS:


- Disk Operating System (DOS)
- DOS/Basic package sold to computer companies
- MS-DOS, including many features from UNIX
- GUI invented in the 1960s, then copied by Apple
- Microsoft copied Apple (Windows working on top of MS-DOS)
- Many OS derived from UNIX (MINIX, LINUX, BSD...)

Device and task oriented OS types:

- Personal Computers (PC)
- Servers: serve users over a network (print, web, IM...) → Solaris, FreeBSD, Linux, Windows Server
- Multi processors: multiple CPU in a single system → Linux, Windows, OS X...
- Handheld computers: PDA, smartphone
- Embedded devices: TV, microwave, DVD player, mp3 player, old cell phones → everything stored in ROM, much more simple OS

More device and task oriented OS types:

- Real-Time: time is key parameter (e.g. assembly line, army, avionics...) → overlap with embedded/handheld systems
- Mainframe: room-sized computers, data centers → OS oriented toward processing many jobs at once and efficient I/O
- Sensor node: tiny computers communicating between each other and a base station (guard border, intrusion/fire detection etc...). Composed of CPU RAM ROM (+other sensors), small battery → simple OS design TinyOS
- Smart card: credit card size with a CPU chip, severe memory/processing constraints → smallest/primitive OS



A computer is often composed of:

- CPU
- Memory
- Monitor + video controller
- Keyboard + keyboard controller
- Hard Disk Drive (HDD) + hard disk controller
- Bus

What are the controllers, and the bus?

## Basics:

- CPU is the “computer’s brain”
- CPU can only execute a specific set of instructions
- CPU fetches instructions from the memory and executes them

## Registers:

- General register: hold variables/temporary results  
e.g. program counter: address of next instruction to fetch
- Stack pointer: parameters/variables not kept in registers
- Program Status Word (PSW): control bits

# Pipeline vs. superscalar





## Superscalar:

- Multiple execution units  
e.g. one for float, int, boolean
- Multiple instructions fetched and decoded at a time
- Instructions held in buffer to be executed
- Issue: no specific order to execute buffered instructions

## Pipeline:

- Execute instruction  $n$ , decode  $n + 1$  and fetch  $n + 2$
- Any fetched instruction must be executed
- Issue: conditional statements





CPUs and cores:

- A CPU core can hold the state of more than one thread
- A core can switch between threads in a nanosecond time scale
- The OS sees several CPUs instead of one core
- Issue: what happens if there are more than two such cores?



Modern terminology<sup>1</sup>:

- CPU: computing component (the physical entity)
- Number of cores: number of independent CPUs in a computing component
- Number of threads: maximum number of instructions that can be passed through or processed by a single core
- Number of logical cores: number of cores times number of threads

---

<sup>1</sup>source: **ARK**

## Access time


1 ns

2 ns

10 ns

10 ms

10 s



## Capacity

&lt; 1 KB

4 MB

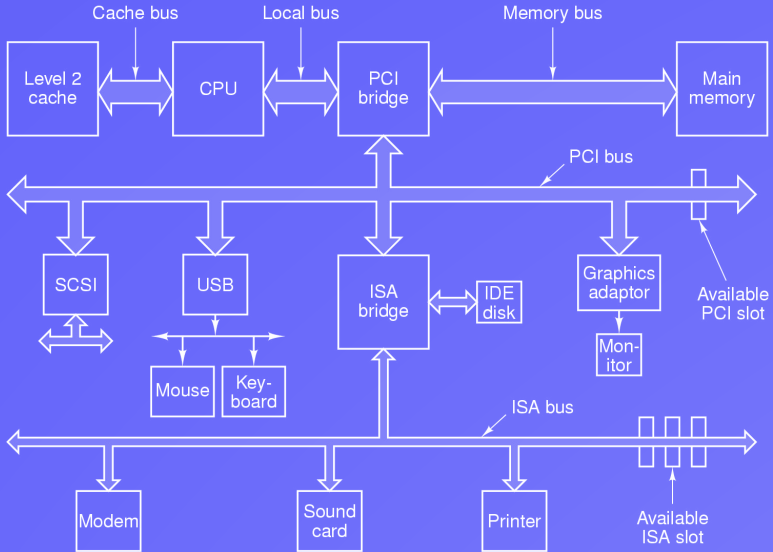
1–8 GB


200–1000 GB

400–800 GB

## Memory types:

- Random Access Memory (RAM): volatile
- Read Only Memory (ROM)
- Electrically Erasable PROM (EEPROM) and flash memory: slower than RAM, non volatile.
- CMOS: save time and date , BIOS parameters
- HDD: divided into cylinder, track and sector






Five major components of an OS:

- Input/Output
- Protection/Security
- Processes
- File system
- System calls


Using the device controller:

- I/O devices and CPU can execute concurrently
- Device controller in charge of a particular device
- Drive controllers have a buffer
- Buffer used to move data to/from buffer to/from memory
- Device controllers interact with the CPU using interrupts



Starting an I/O device and getting an interrupt:

- 1 Send instruction to controller
- 2 Controller signals the end to interrupt controller
- 3 Assert pin to interrupt the CPU
- 4 Send extra information



Starting an I/O device and getting an interrupt:

- 1 Send instruction to controller
- 2 Controller signals the end to interrupt controller
- 3 Assert pin to interrupt the CPU
- 4 Send extra information

Processing an interrupt:

- 1 Push user program counter and PSW onto the stack
- 2 Switch to kernel mode
- 3 Find the device interrupt handler in the interrupt vector
- 4 Query the device for its status
- 5 Return to user program (1st instruction not yet executed)



Remarks of interrupts:

- Incoming interrupts are disabled during the process
- Software can generate interrupts: trap  
e.g. java exception, division by 0...
- An OS is almost always interrupt driven

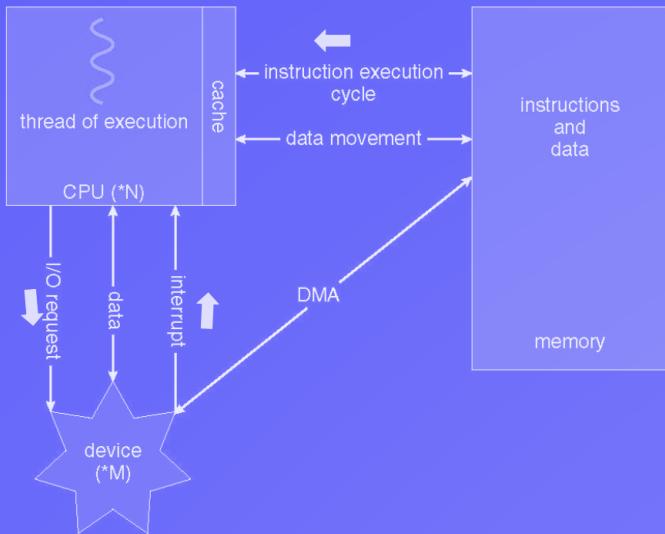
Simplest method:

- 1 Call the driver
- 2 Start the I/O
- 3 Wait in a tight loop
- 4 Continuously poll the device to know its state

Disadvantage: CPU *busy waiting* until device is finished

### Direct Memory Access (DMA):

- High speed I/O devices
- Transmit information close to memory speeds
- Device controller directly transfers the blocks of data from the buffer to the main memory
- No help from the CPU
- One interrupt per block, instead of one per byte



## CPU:

- Kernel Mode:
  - Set using a bit in the PSW
  - Any CPU instruction and hardware feature are available
- User mode:
  - Only a subset of instructions/features is available
  - Setting PSW kernel mode bit forbidden

## Memory:

- Base and limit registers: holds smallest legal physical memory address and size of range, respectively
- Memory outside the address space is protected

## Input/Output:

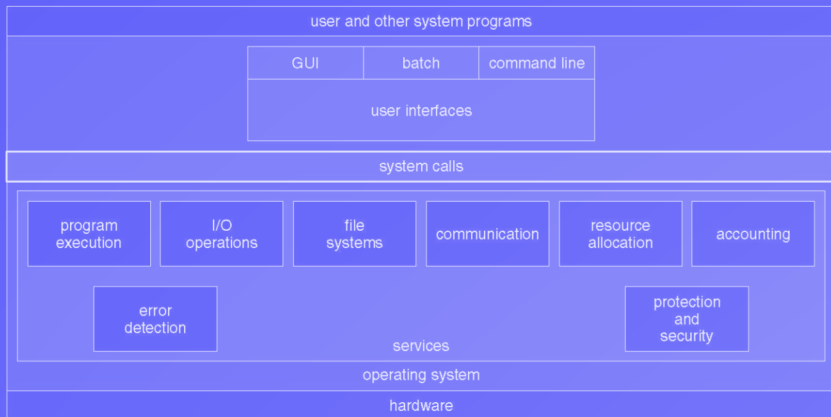
- All I/O instructions are privilege instructions
- OS treats I/O operations to ensure correctness and legality

A process holds all the necessary information to run a program:

- Address space belonging to the process and containing:
  - Executable program
  - Program's data
  - Program's stack
- Set of resources:
  - Registers
  - List of open files
  - Alarms
  - List of related processes
  - Any other information required by the program

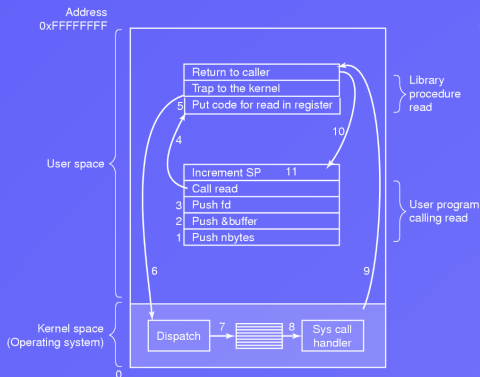
The OS hides peculiarities of the disk and other I/O devices

- Data stored in files
- Files are grouped inside directories
- Top directory is called root directory
- Any file can be specified using its path name
- Each process has a working directory
- Before reading/writing in a file permissions are checked
- If access is granted a file descriptor is returned
- Removable devices can be mounted on the main tree
- Block special files: represent devices such as disks
- Character special files: represent devices that accept or output character stream
- Pipe: pseudo file used to connect two processes





```
ssize_t read(int fd, void *buf, size_t count);
```



Steps for a simple read:

- 1-3: push parameters on the stack
- 5-6: switch to kernel mode
- 7: use a table of pointers to system calls handler to dispatch to the correct one
- 11: increment stack pointer to remove parameters pushed (1-3)

## Partial list of common Unix system calls:

- Processes:

```
pid=fork(); pid=waitpid(pid, &statloc, options);  
s=execve(name, argv, environp); exit(status);
```

- Files:

```
fd=open(file,how,...); s=close(fd); s=stat(name,*buf);  
n=read(fd,buffer,nbytes); n=write(fd,buffer,nbytes);  
position=lseek(fd,offset,whence);
```


- Directory and file system:

```
s=mkdir(name,mode); s=rmdir(name); umount(name); s=unlink(name);  
mount(special,name,flags,types,args); s=link(name1,name2);
```


- Misc:

```
s=chdir(dirname); s=chmod(name,mode); sec=time(*t);  
s=kill(pid,signal);
```

## Monolithic kernel



## Micro kernel







Thank you!