




# Introduction to Operating Systems

## 2. Processes and threads


Manuel – Fall 2019





A *process* is an abstraction of a running program:

- At the core of the OS
- Process is the unit for resource management
- Oldest and most important concept
- Turn a single CPU into multiple virtual CPUs
- CPU quickly switches from process to process
- Each process run for 10-100 ms
- Processes hide the effect of interrupts

One program counter



Four program counters



Multiprogramming strategies and issues:

- CPU switches rapidly back and forth among all the processes
- Rate of computation of a process is not uniform/reproducible
- Potential issue under time constraints; e.g.
  - Read from tape
  - Idle loop for tape to get up to speed
  - Switch to another process
  - Switch back... too late

Differences between programs and processes:

- Running twice a program generates two processes
- Program: sequence of operations to perform
- Process: program, input, output, state


Differences between programs and processes:

- Running twice a program generates two processes
- Program: sequence of operations to perform
- Process: program, input, output, state

Example.

Describe the process of baking a cake when the phone rings...

Base+limit



Four main events causing process to be created:

- System initialization
- Execution of a “process creation” system call
- A user requests a new process
- Initiation of a batch job



Unix like systems:

- Creating a new process is done using one system call: `fork()`
- The call creates an exact clone of the calling process
- Child process executes `execve` to run a new program

Windows system:

- Function call `CreateProcess`, creates a new process and loads the program in the new process
- This call takes 10 parameters


Parent and child have their own address space and a change in one is invisible to the other

Any created processes ends at some stage:

- Normal exit (voluntary)  
work is done, execute a system call to tell OS it is finished  
`exit`, `ExitProcess`
- Error exit (voluntary)  
e.g. requested file does not exist
- Fatal error (involuntary)  
e.g. referencing non existent memory, dividing by 0
- Killed by another process (involuntary)  
`kill`, `TerminateProcess`


Two main approaches:

- UNIX like systems:
  - Parent creates a child
  - Child can create its own child
  - The hierarchy is called *process group*
  - Impossible to disinherit a child
- Windows system:
  - All processes are equal
  - A parent has a token to control its child
  - Token can be given to another process



Possible states:

- 1 Waiting for some input
- 2 Scheduler picks another process
- 3 Scheduler picks this process
- 4 Input becomes available



Possible states:

- ① Waiting for some input
- ② Scheduler picks another process
- ③ Scheduler picks this process
- ④ Input becomes available

Change of perspective on the inside of the OS:

- Do not think in terms of interrupt but of process
- Lowest level of the OS is the scheduler
- Interrupt handling, starting/stopping processes are hidden in the scheduler

A simple model for processes:


- Each process is represented using a structure called *process control block*
- The structure contains important information such as:
  - State
  - Program counter
  - Stack pointer
  - Memory allocation
  - Open files
  - Scheduling information
  - ...
- All the processes are stored in an array called *process table*

| Process management    | Memory management             | File management   |
|-----------------------|-------------------------------|-------------------|
| registers             | pointer to text segment info  | root directory    |
| program counter       | pointer to data segment info  | working directory |
| program status word   | pointer to stack segment info | file descriptors  |
| stack pointer         |                               | user ID           |
| process state         |                               | group ID          |
| priority              |                               |                   |
| scheduling parameters |                               |                   |
| process ID            |                               |                   |
| parent process        |                               |                   |
| process group         |                               |                   |
| signals               |                               |                   |
| starting time         |                               |                   |
| CPU time used         |                               |                   |
| children's CPU time   |                               |                   |
| next alarm            |                               |                   |

Lowest OS level:

- 1 Push user program counter, PSW...on stack
- 2 Hardware loads new program counter from interrupt vector
- 3 Save registers (assembly)
- 4 Setup new stack (assembly)
- 5 Finish up the work for the interrupt
- 6 Scheduler decides which process to run next
- 7 Load and run the “new current process”, i.e. memory map, registers...(assembly)





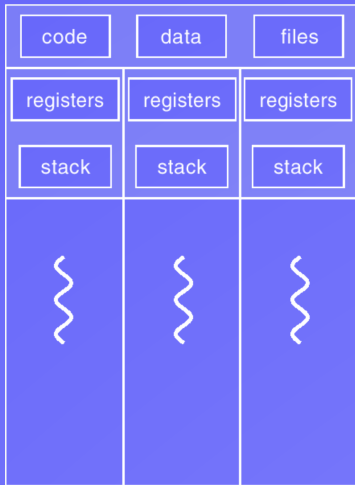

A thread is the basic unit of CPU utilisation consisting of:

- Thread ID
- Program counter
- Register set
- Stack space

All the threads within a process share:

- Code section
- Data section
- OS resources


# Single vs. multi-threaded



## Processes and threads:

- A thread has the same possible states as a process
- Transitions are similar to the case of a process
- Threads are sometimes called lightweight process
- No protection is required for threads, compared to processes
- A process starts with one threads and can create more
- Processes want as much CPU as they can
- Threads can give up the CPU to let others using it

|   |   |  |  |   |   |
|---|---|--|--|---|---|
| <p>Three score and seven years ago, our fathers brought forth upon this continent a new nation conceived in liberty and dedicated to the proposition that all men are created equal. Now we are engaged in a great civil war testing whether that</p> | <p>nation, in any nation so conceived and so dedicated, can long endure. We are currently in a great battlefield of the war.</p> <p>We have come to dedicate a portion of that field as a final resting place for those who here gave their</p> | <p>lives that this nation might live. It is altogether fitting and proper that we should do this.</p> <p>But in a larger sense, we cannot dedicate, we cannot consecrate we cannot hallow this ground. The brave men, living and dead,</p> | <p>who struggled here have consecrated it, far above our poor power to add or detract. The world will little note, nor long remember, what we say here, but it can never forget what they did here. It is for us the living, rather, to be dedicated</p> | <p>here to the unfinished work which they who fought here have devoted to us. It is rather for us to be here dedicated to the great task remaining before us, that these brave honest dead will have increased devotion to that cause for which</p> | <p>they gave the last full measure of devotion, that we here highly resolve that these dead shall not have died in vain that this nation, under God, shall have a new birth of freedom and that the government of the people, by the people, for the people</p> |
|---|---|--|--|---|---|




Thread share many data structure:

- A thread could close a file that another thread is reading
- A thread notices a lack of memory and allocate more. A thread switch occurs, the new threads also notices the lack of memory and also allocates some

Should a child inherit all the threads form its parents?

- No: the child might not function properly
- Yes: if a parent thread was waiting for some keyboard input, who gets it?



The pthread library has over 60 function calls, important ones are:

- Create a thread:

```
int pthread_create(pthread_t *thread, const pthread_attr_t
*attr,void *(*start_routine) (void *), void *arg);
```

- Terminate a thread:

```
void pthread_exit(void *retval);
```

- Wait for a specific thread to end:

```
int pthread_join(pthread_t thread, void **retval);
```

- Release CPU to let another thread run:

```
int pthread_yield(void);
```

- Create and initialise a thread attribute structure:

```
int pthread_attr_init(pthread_attr_t *attr);
```

- Delete a thread attribute object:

```
int pthread_attr_destroy(pthread_attr_t *attr);
```




Write a program that creates 10 threads, and prints their ID.

## threads.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <pthread.h>
4  #define THREADS 10
5  void *gm(void *tid) {
6      printf("Good morning from thread %lu\n",*(unsigned long int*)tid);
7      pthread_exit(NULL);
8  }
9  int main () {
10     int status, i; pthread_t threads[THREADS];
11     for(i=0;i< THREADS;i++) {
12         printf("thread %d\n",i);
13         status=pthread_create(&threads[i],NULL,gm,(void*)&(threads[i]));
14         if(status!=0) {
15             fprintf(stderr,"thread %d failed with error %d\n",i,status);
16             exit(-1);
17         }
18     }
19     exit(0);
20 }
```

# Threads in user-space – N:1




## User-space threads:

- Kernel thinks it manages single threaded processes
- Threads implemented in a library
- Thread table similar to process table, managed by run-time system
- Switching thread does not require to trap the kernel

## Questions.

- What if a thread issues a blocking system call?
- Threads within a process have to voluntarily give up the CPU

# Thread in the kernel – 1:1




## Kernel space thread:

- Kernel manages the thread table
- Kernel calls are issued to request a new thread
- Calls that might block a thread are implemented as system call
- Kernel can run another thread in the meantime

## Questions.

- Why does it have a much higher cost than user space threads?
- Signals are sent to processes, which thread should received it?

# Hybrid threads – M:N



## Hybrid threads:

- Compromise between user-level and kernel-level
- Threading library schedules user threads on available kernel threads

## Questions.

- How to implement hybrid threads?
- How to handle scheduling?

Best thread approach:

- Hybrid looked attractive
- Most systems are coming back to 1:1
- Different approaches exist on how to use threads  
e.g. thread blocks on “receive system call” vs. pop up threads
- Switching implementation from single thread to multiple thread is not easy task
- Requires redesigning the whole system
- Backward compatibility must be preserved
- Research still going on to find better ways to handle threads





Thank you!